

# CS 51 Project: Scalable Bloom Filters, the Algorithms

May 1, 2015

Joseph Kahn, Grace Lin, Aron Szanto  
Cambridge, MA 02138

## Contents

<b>1 Static Bloom Filters</b>	<b>1</b>
<b>2 Scalable Bloom Filters</b>	<b>2</b>
<b>3 Partitioned Bloom Filters</b>	<b>3</b>

## 1 Static Bloom Filters

Fundamentally, a Bloom filter is a probabilistic data structure. By giving up the ability to definitively check the presence of an element in the data structure, a Bloom filter is able to store elements with a constant element/memory ratio, regardless of the size of the element. The size of the bitset,  $m$ , and the number of hash functions,  $k$ , can be modified to accomodate any number of inputs,  $n$ , with any desired false positive percentage, *errbound*.

The probability for a given bit being set by a particular hash function for a particular entry is  $\frac{1}{m}$ . Therefore, the probability for the bit of interest not to be set is  $1 - \frac{1}{m}$ . It follows that the probability for a particular bit to still be 0 after  $n$  inputs is  $(1 - \frac{1}{m})^{nk}$ . Therefore, the probability for the bit to be set to 1 after  $n$  inputs is  $1 - (1 - \frac{1}{m})^{nk}$ , which for very large  $n$  and  $m$  can be approximated as  $(1 - e^{-\frac{kn}{m}})^k$ .

Note: This approach assumes that the hash functions are independent and that there will be no collisions across hash functions. However, for very large  $m, n$ , this approximation is close enough to calculate a false positive percentage.

As derived above,  $errbound = (1 - e^{-\frac{kn}{m}})^k$ . For a given  $m, n$ , we can treat  $errbound$  as a single-variable function of  $k$ , and using some simple calculus, find the value of  $k$  with respect to at the minimum  $errbound$ . The zero of the derivative of  $errbound$  is  $k = \frac{m}{n} \ln(2)$ . Substituting this  $k$  value back into  $errbound$ ,  $errbound$  simplifies to  $errbound = (1 - e^{-\ln(2)})^{\frac{m}{n} \ln(2)}$ , which cleans up into  $\ln(errbound) = -\frac{m}{n} (\ln(2))^2$ . Since  $errbound$  and  $n$  are given, we can thus derive  $m = -\frac{n \ln(errbound)}{(\ln(2))^2}$ . With this derivation of  $m$ , we can also simplify  $k$  to be  $-\frac{\ln(errbound)}{\ln 2}$ .

The explanation for the optimal  $fill$  value is as follows: given an error bound and an expected number of elements, the optimal  $k$  value is  $k = \frac{m}{n} \ln 2$ .  $fill = 1 - e^{-\frac{km}{n}}$ . Substituting  $\ln 2 = \frac{km}{n}$ ,  $fill = \frac{1}{2}$ .

## 2 Scalable Bloom Filters

A scalable Bloom filter (SBF) deals with the problem of having to a priori declare the expected number of inputted elements when creating a Bloom filter. While a Bloom filter will never reject an input, if the number of bits set to 1, or the fill ratio, increases beyond a certain threshold (we show above that this threshold is  $\frac{1}{2}$ ), then the false positive probability of the Bloom filter will increase beyond the desired error bound. We can overcome this problem by creating a series of Bloom filters. The initial Bloom filter is simply a static Filter of some arbitrary size, and once this Bloom filter has been filled to a certain fill ratio, a subsequent Bloom filter is created, and so on. The SBF struct stores several things: an array of all the existing Bloom filters, the number of Bloom filters, a scaling factor that determines how much bigger each subsequent Bloom filter will be, and a tightening ratio that determines how much smaller the error bound on each subsequent Bloom filter will be.

We will show that such error bounds will converge. For a SBF with  $l$  filters, initial error bound  $e_0$  and tightening ratio  $r$ , with  $e_i = e_0 * r^i$ . The overall errorbound,  $error$ , is equal to  $1 - \prod_{i=0}^N (1 - e_i)$ . This is bounded by  $\prod_{i=0}^N e_i = \frac{e_0}{1-r}$ . Thus,  $error$  converges to approximately  $\frac{e_0}{1-r}$ . Thus, given a particular total error bound and an  $r$  value, it is possible to calculate the necessary error bound for each static filter within the SBF.

### 3 Partitioned Bloom Filters

For the third part of our project, we implemented partitioned Bloom filters in scalable Bloom filters. Partitioned Bloom filters are more efficient because they prevent collisions/overlaps between hash functions. Algorithmically, calculating fill ratios and error bounds for each partition are equivalent to setting  $k$  to 1. The total error bound is  $err_{tot} = fill^k$ , where  $fill$  is the fill ratio of a particular slice ( $e = 1 - (1 - \frac{1}{m})^n$ , where  $m$  is the size of the slice. Using a Taylor expansion,  $fill$  can be approximated by  $1 - e^{-\frac{n}{m}}$ . It easily follows, through some algebra, that  $n$  can be approximated by  $M \frac{\ln(1-fill)}{-k}$ . This gives us a  $M$  value, which is the size of the bitset, of  $n \frac{-k}{1-fill}$ .  $fill^k = err_{tot}$ , so with a fill of  $\frac{1}{2}$ , we get  $k = -\frac{\ln(err_{tot})}{\ln 2}$ . So,  $M = n \frac{2 \ln err_{tot}}{\ln(2)}$ . If we compare the coefficient of  $M$  and the coefficient of  $m$  from the static classic Bloom filter previously discussed, the coefficient for  $M$  is about 20 percent smaller than that of  $m$ . Given the same number of values and the same total error bound, a partitioned Bloom filter is about 20 percent more memory efficient than a standard Bloom filter. Comparatively, had we chosen to increase  $k$  significantly rather than partition the bitset, the number of collisions would have also increased significantly and thus reduced the efficiency of our Bloom filter.